

# Nagilli Chitti Keerthana

Jersey City, NJ (Open to Relocation) | [nckeerthana120@gmail.com](mailto:nckeerthana120@gmail.com) | +1 (551)-358-7650 | [LinkedIn](#) | [GitHub](#)

## Summary

AI engineer skilled in building full-stack LLM-powered systems and API-based integrations, with experience in multi-agent orchestration, structured data extraction, and production AI deployment on Azure. 2 published papers. Authorized to work in the US.

## Education

### Master of Science in Applied Artificial Intelligence

Aug 2024 – May 2026

Stevens Institute of Technology, Hoboken, NJ, USA

GPA: 3.78/4.00 | Relevant Coursework: Applied Machine Learning, Natural Language Processing, Probability & Statistics

### Bachelor of Engineering in Electronics and Communication

Aug 2020 – May 2024

PES University, Bengaluru, Karnataka, India

## Technical Skills

**AI & ML:** LLMs, LangGraph, LangChain, RAG, Prompt Engineering, Sentence Transformers, PyTorch, TensorFlow, NLP

**Dev:** Python, C/C++, FastAPI, Flask, TypeScript, JavaScript (React, Next.js), Docker, Git, SQL, Evaluation

**Cloud & Data:** Azure, PostgreSQL, Snowflake, Apache Kafka, Airflow, DBT, Pandas

## Experience

### Stevens Institute of Technology, Graduate Research Assistant – Hoboken, NJ, USA

Jan 2025 – Dec 2025

- Designed a TensorFlow-based multi-agent reinforcement learning (MADDPG) model with custom reward engineering and actor-critic optimization, improving the target objective by 76% over a mathematical baseline.

### NxtConnect AI, AI Intern – NYC, New York, USA

Jun 2025 – Aug 2025

- Integrated Claude AI via REST APIs with prompt engineering to extract structured entities from customer conversations, building a Neo4j-based GraphRAG pipeline for a conversational CRM chatbot that achieved 90% retrieval accuracy.
- Built an automated evaluation pipeline, deployed via Docker, that scored PyTorch BERT event detection outputs against a labeled ground truth set, flagging low confidence cases for human review to validate 85.7% accuracy.
- Implemented a multimodal AI pipeline with Whisper and EasyOCR in a FastAPI backend and TypeScript frontend, enabling in-chat processing of audio, PDFs, images, and videos with 40× faster processing and 90% video-text accuracy.

### Shadowfax Pvt Ltd, Data Analyst Intern – Bengaluru, India

Jan 2024 – Apr 2024

- Streamlined ETL pipelines and weekly performance reporting using Python, SQL, and the Google Sheets API, analyzing 580k+ member engagement records across 1,100+ hubs to deliver real-time insights while reducing manual effort.

### Nokia Bell Labs, Research Assistant – Bengaluru, India

Aug 2023 – Jan 2024

- Evaluated predictive models for network handover success, benchmarking XGBoost and Random Forest against Graph Neural Networks, achieving 25% improvement in prediction accuracy.

## Projects

### Audible Frames on Azure Cloud

May 2026 – Jun 2026

Azure AI Vision, Azure OpenAI, Azure AI Speech, Content Safety, FastAPI, Docker

[GitHub](#)

- Designed and deployed an image-to-audio pipeline on Azure (AI Vision, OpenAI Responses API, Azure Speech) via GitHub Actions CI/CD, using Insights dashboards for monitoring while improving caption quality 38% through A/B testing.

### Multi-Agent AI System with LangGraph

Jan 2026 – May 2026

LangGraph, MCP (Tavily), CLIP, FAISS, GPT-4V, Python

[GitHub](#)

- Automated report generation and PDF query response by building a 4-agent LangGraph system with MCP tool calling and a CLIP and FAISS RAG pipeline enabling GPT-4V to retrieve text and images.

## Research & Publications

- Best Research Paper Award** – 12th International Conference on Recent Challenges in Engineering and Technology for “Power Load Demand Prediction using LSTM and Actor-Critic Model” among 269 papers.
- First author** of “MADDPG-Driven Uplink Secrecy Optimization for Multi-User MIMO-NOMA Networks” published at ICCIT 2025, received **Second Prize** at Stevens Institute of Technology.